

```
In [ ]: from sklearn.datasets import load_diabetes
import pandas as pd

diabetes = load_diabetes()
```

```
In [ ]: diabetes
```

```

Out[ ]: {'data': array([[ 0.03807591,  0.05068012,  0.06169621, ..., -0.00259226,
    0.01990749, -0.01764613],
  [-0.00188202, -0.04464164, -0.05147406, ..., -0.03949338,
    -0.06833155, -0.09220405],
  [ 0.08529891,  0.05068012,  0.04445121, ..., -0.00259226,
    0.00286131, -0.02593034],
  ...,
  [ 0.04170844,  0.05068012, -0.01590626, ..., -0.01107952,
    -0.04688253,  0.01549073],
  [-0.04547248, -0.04464164,  0.03906215, ...,  0.02655962,
    0.04452873, -0.02593034],
  [-0.04547248, -0.04464164, -0.0730303 , ..., -0.03949338,
    -0.00422151,  0.00306441]]),
  'target': array([151.,  75., 141., 206., 135.,  97., 138.,  63., 110., 310., 101.,
    69., 179., 185., 118., 171., 166., 144.,  97., 168.,  68.,  49.,
    68., 245., 184., 202., 137.,  85., 131., 283., 129.,  59., 341.,
    87.,  65., 102., 265., 276., 252.,  90., 100.,  55.,  61.,  92.,
   259.,  53., 190., 142.,  75., 142., 155., 225.,  59., 104., 182.,
   128.,  52.,  37., 170., 170.,  61., 144.,  52., 128.,  71., 163.,
   150.,  97., 160., 178.,  48., 270., 202., 111.,  85.,  42., 170.,
   200., 252., 113., 143.,  51.,  52., 210.,  65., 141.,  55., 134.,
    42., 111.,  98., 164.,  48.,  96.,  90., 162., 150., 279.,  92.,
    83., 128., 102., 302., 198.,  95.,  53., 134., 144., 232.,  81.,
   104.,  59., 246., 297., 258., 229., 275., 281., 179., 200., 200.,
   173., 180.,  84., 121., 161.,  99., 109., 115., 268., 274., 158.,
   107.,  83., 103., 272.,  85., 280., 336., 281., 118., 317., 235.,
    60., 174., 259., 178., 128.,  96., 126., 288.,  88., 292.,  71.,
   197., 186.,  25.,  84.,  96., 195.,  53., 217., 172., 131., 214.,
    59.,  70., 220., 268., 152.,  47.,  74., 295., 101., 151., 127.,
   237., 225.,  81., 151., 107.,  64., 138., 185., 265., 101., 137.,
   143., 141.,  79., 292., 178.,  91., 116.,  86., 122.,  72., 129.,
   142.,  90., 158.,  39., 196., 222., 277.,  99., 196., 202., 155.,
    77., 191.,  70.,  73.,  49.,  65., 263., 248., 296., 214., 185.,
    78.,  93., 252., 150.,  77., 208.,  77., 108., 160.,  53., 220.,
   154., 259.,  90., 246., 124.,  67.,  72., 257., 262., 275., 177.,
    71.,  47., 187., 125.,  78.,  51., 258., 215., 303., 243.,  91.,
   150., 310., 153., 346.,  63.,  89.,  50.,  39., 103., 308., 116.,
   145.,  74.,  45., 115., 264.,  87., 202., 127., 182., 241.,  66.,
    94., 283.,  64., 102., 200., 265.,  94., 230., 181., 156., 233.,
    60., 219.,  80.,  68., 332., 248.,  84., 200.,  55.,  85.,  89.,
    31., 129.,  83., 275.,  65., 198., 236., 253., 124.,  44., 172.,
   114., 142., 109., 180., 144., 163., 147.,  97., 220., 190., 109.,
   191., 122., 230., 242., 248., 249., 192., 131., 237.,  78., 135.,
   244., 199., 270., 164.,  72.,  96., 306.,  91., 214.,  95., 216.,
   263., 178., 113., 200., 139., 139.,  88., 148.,  88., 243.,  71.,
    77., 109., 272.,  60.,  54., 221.,  90., 311., 281., 182., 321.,
    58., 262., 206., 233., 242., 123., 167.,  63., 197.,  71., 168.,
   140., 217., 121., 235., 245.,  40.,  52., 104., 132.,  88.,  69.,
   219.,  72., 201., 110.,  51., 277.,  63., 118.,  69., 273., 258.,
    43., 198., 242., 232., 175.,  93., 168., 275., 293., 281.,  72.,
   140., 189., 181., 209., 136., 261., 113., 131., 174., 257.,  55.,
    84.,  42., 146., 212., 233.,  91., 111., 152., 120.,  67., 310.,
    94., 183.,  66., 173.,  72.,  49.,  64.,  48., 178., 104., 132.,
   220.,  57.]),
  'frame': None,
  'DESCR': '.._diabetes_dataset:\n\nDiabetes dataset\n-----\n\nTen baselin
e variables, age, sex, body mass index, average blood\npressure, and six blood serum
measurements were obtained for each of n=\n442 diabetes patients, as well as the res
ponse of interest, a\nquantitative measure of disease progression one year after base
line.\n\n**Data Set Characteristics:**\n\n :Number of Instances: 442\n\n :Number of

```

Attributes: First 10 columns are numeric predictive values\n\n :Target: Column 11 is a quantitative measure of disease progression one year after baseline\n\n :Attribute Information:\n - age age in years\n - sex\n - bmi body mass index\n - bp average blood pressure\n - s1 tc, total serum cholesterol\n - s2 ldl, low-density lipoproteins\n - s3 hdl, high-density lipoproteins\n - s4 tch, total cholesterol / HDL\n - s5 ltcg, possibly log of serum triglycerides level\n - s6 glu, blood sugar level\n\nNote: Each of these 10 feature variables have been mean centered and scaled by the standard deviation times the square root of `n_samples` (i.e. the sum of squares of each column totals 1).\n\nSource URL:\n<https://www4.stat.ncsu.edu/~boos/var.select/diabetes.html>\n\nFor more information see:\nBradley Efron, Trevor Hastie, Iain Johnstone and Robert Tibshirani (2004) "Least Angle Regression," Annals of Statistics (with discussion), 407-499.\n(https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle_2002.pdf)\n'

```
'feature_names': ['age',
                  'sex',
                  'bmi',
                  'bp',
                  's1',
                  's2',
                  's3',
                  's4',
                  's5',
                  's6'],
'data_filename': 'diabetes_data_raw.csv.gz',
'target_filename': 'diabetes_target.csv.gz',
'data_module': 'sklearn.datasets.data'}
```

```
In [ ]: dir(diabetes)
```

```
Out[ ]: ['DESCR',
         'data',
         'data_filename',
         'data_module',
         'feature_names',
         'frame',
         'target',
         'target_filename']
```

```
In [ ]: len(diabetes.target)
```

```
Out[ ]: 442
```

```
In [ ]: diabetesdf = pd.DataFrame(data=diabetes.data,
                                  columns=diabetes.feature_names)
```

```
In [ ]: diabetesdf
```

Out[]:

	age	sex	bmi	bp	s1	s2	s3	s4	s5
0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	-0.002592	0.019907
1	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	-0.039493	-0.068332
2	0.085299	0.050680	0.044451	-0.005670	-0.045599	-0.034194	-0.032356	-0.002592	0.002861
3	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	0.034309	0.022688
4	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	-0.002592	-0.031988
...
437	0.041708	0.050680	0.019662	0.059744	-0.005697	-0.002566	-0.028674	-0.002592	0.031193
438	-0.005515	0.050680	-0.015906	-0.067642	0.049341	0.079165	-0.028674	0.034309	-0.018114
439	0.041708	0.050680	-0.015906	0.017293	-0.037344	-0.013840	-0.024993	-0.011080	-0.046883
440	-0.045472	-0.044642	0.039062	0.001215	0.016318	0.015283	-0.028674	0.026560	0.044529
441	-0.045472	-0.044642	-0.073030	-0.081413	0.083740	0.027809	0.173816	-0.039493	-0.004222

442 rows × 10 columns

```

In [ ]: from sklearn.svm import SVR
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler

epsilon=0.2

svr_rbf = SVR(kernel="rbf", C=100, gamma=0.1, epsilon=epsilon)
svr_lin = SVR(kernel="linear", C=100, gamma="auto")
svr_poly = SVR(kernel="poly", C=100, gamma="auto", degree=3, epsilon=epsilon, coef0=1)

```

This data looks treated

```

In [ ]: X= diabetesdf

scaler = StandardScaler()
X = scaler.fit_transform(X)

```

```

In [ ]: X

```

```
Out[ ]: array([[ 0.80050009,  1.06548848,  1.29708846, ..., -0.05449919,
           0.41853093, -0.37098854],
          [-0.03956713, -0.93853666, -1.08218016, ..., -0.83030083,
           -1.43658851, -1.93847913],
          [ 1.79330681,  1.06548848,  0.93453324, ..., -0.05449919,
           0.06015558, -0.54515416],
          ...,
          [ 0.87686984,  1.06548848, -0.33441002, ..., -0.23293356,
           -0.98564884,  0.32567395],
          [-0.9560041 , -0.93853666,  0.82123474, ...,  0.55838411,
           0.93616291, -0.54515416],
          [-0.9560041 , -0.93853666, -1.53537419, ..., -0.83030083,
           -0.08875225,  0.06442552]])
```

```
In [ ]: y= diabetes.target
```

```
In [ ]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=
```

```
In [ ]: svr_rbf.fit(X_train, y_train)
svr_lin.fit(X_train, y_train)
svr_poly.fit(X_train, y_train)
```

```
Out[ ]: SVR
SVR(C=100, coef0=1, epsilon=0.2, gamma='auto', kernel='poly')
```

```
In [ ]: y_pred1 = svr_rbf.predict(X_test)
y_pred2 = svr_lin.predict(X_test)
y_pred3 = svr_poly.predict(X_test)
```

```
In [ ]: len(y_pred3)
```

```
Out[ ]: 89
```

```
In [ ]: from sklearn.metrics import mean_squared_error

mean_squared_error(y_test, y_pred1)
```

```
Out[ ]: 2601.7743632899696
```

```
In [ ]: from sklearn.metrics import mean_squared_error

mean_squared_error(y_test, y_pred2)
```

```
Out[ ]: 3007.3560832247504
```

```
In [ ]: from sklearn.metrics import mean_squared_error

mean_squared_error(y_test, y_pred3)
```

```
Out[ ]: 3749.257378564513
```

svr_rbf.predict is the best

In []: